

Companion Notes for `landau_hermite_jax.py`: Fast Landau–Hermite Collisions via SOE→MPO/TT (JAX-first)

(auto-generated companion for the standalone script)

February 3, 2026

Abstract

This document explains the numerics, implementation, and test suite used by the standalone script `landau_hermite_jax.py`. The script implements a fast evaluation of the spatially homogeneous Landau (Fokker–Planck) collision operator in a 3D tensor-product Hermite basis, using a separable quadrature (“SOE” style separation), tensor-product operator factorizations (“MPO” viewpoint), and optional tensor-train (TT) compression (currently used in the NumPy debug path; the JAX path relies on JIT+vectorized contractions).

The goal is pedagogic completeness: what the code computes, how it is normalized, how it is time-stepped, what “linearized” means in this context, why the tests are designed the way they are, and how to interpret the produced figures (Fig. 1/Fig. 2 panels and the test plots).

Contents

1 Quickstart and reproducibility	3
1.1 What the script produces	3
1.2 Backend	3
1.3 CLI options	3
2 Physical model: Landau collision operator	5
2.1 Continuous operator (context)	5
2.2 What the script discretizes	5
3 Normalization, basis, and coefficient tensor representation	6
3.1 Thermal normalization	6
3.2 Hermite basis used in the code	6
3.3 Tensor coefficient representation	6
3.4 Moment diagnostics and invariants	6
4 Fast collision evaluation: SOE→MPO/TT	7
4.1 Why naive evaluation is expensive	7
4.2 SOE / separable quadrature idea	7
4.3 Separable Hankel structure	7
4.4 MPO viewpoint: tensor-product operator contraction	8
4.5 TT (tensor train) compression	8
4.6 Algorithmic walkthrough (as implemented)	8

5	Nonlinear vs linearized evolution	9
5.1	Nonlinear operator	9
5.2	Linearization about a Maxwellian	9
5.3	Entropy vs free energy for linearized dynamics	9
6	Time stepping and JAX implementation	10
6.1	Integrator	10
6.2	JAX performance strategy	10
7	Initial conditions used in Fig. 1 and Fig. 2	10
7.1	Fig. 1: one-species IC (two-stream-like, positivity-safe)	10
7.2	Fig. 1: two-species IC (positivity-safe hot/equilibrium)	10
7.3	Fig. 2: far-from-Maxwellian ICs	10
8	Figures produced by the script	11
8.1	Main panels	11
8.2	Test suite summary and plots	11
9	Interpreting results and limitations	19
9.1	Positivity	19
9.2	Monotonicity	19
9.3	Truncation and parameter choices	19
10	References (selected)	19

1 Quickstart and reproducibility

1.1 What the script produces

Running `python landau_hermite_jax.py` (defaults) writes:

- `Fig1_panel.pdf` and `Fig1_panel.png`: a 2×2 diagnostic panel for a one-species relaxation problem and a two-species equilibration problem.
- `Fig2_panel.pdf` and `Fig2_panel.png`: same dynamics/parameters as Fig. 1, but with a different, strongly non-Maxwellian initial condition (IC) designed to remain nonnegative on diagnostic grids.
- Optionally (if `--run_tests`), a timestamped directory under `tests_landau_hermite/run_YYYYmmdd_HHMMSS/` containing plots, CSV, and JSON summaries. The latest run is also linked as `tests_landau_hermite/latest`.

1.2 Backend

The script has two backends:

- **JAX backend** (`--backend jax`, default): uses `jax.jit`, `vmap/einsum`, and `lax.scan` time stepping for performance and scalability.
- **NumPy backend** (`--backend numpy`): for debugging and for certain explicit checks.

Both backends use the *same* precomputed coefficient tables and the same fast SOE→MPO contraction structure; the difference is whether contractions are executed with NumPy or JAX/XLA.

1.3 CLI options

The script self-documents its CLI via `-h`. For convenience we include the exact help text:

```
usage: landau_hermite_jax.py [-h] [--backend {jax, numpy}] [--nmax NMAX]
                               [--Q Q] [--maxK MAXK]
                               [--integrator {rk2, ssprk3, rk4}]
                               [--linearized {on, off}]
                               [--linearized_method {tangent, matrix}]
                               [--progress_chunks PROGRESS_CHUNKS] [--quiet]
                               [--use_tt] [--tt_tol TT_TOL] [--tt_rmax TT_RMAX]
                               [--dt_1sp DT_1SP] [--dt_2sp DT_2SP]
                               [--tmax_1sp TMAX_1SP] [--tmax_2sp TMAX_2SP]
                               [--steps_1sp STEPS_1SP] [--steps_2sp STEPS_2SP]
                               [--fig1_ic {prl_m2, twoStream}]
                               [--no_enforce_nonneg_ic] [--amp1 AMP1]
                               [--dT2 DT2] [--outprefix_fig1 OUTPREFIX FIG1]
                               [--outprefix_fig2 OUTPREFIX FIG2]
                               [--outprefix OUTPREFIX] [--skip_fig1]
                               [--skip_fig2] [--fig2_strength FIG2_STRENGTH]
                               [--seed SEED] [--entropy_nx ENTROPY_NX]
                               [--entropy_xlim ENTROPY_XLIM] [--run_tests]
                               [--tests_outdir TESTS_OUTDIR]
                               [--tests_nmax_list TESTS_NMAX_LIST]
                               [--tests_reps_rhs TESTS_REPS_RHS]
                               [--tests_reps_bench TESTS_REPS_BENCH]
                               [--tests_steps TESTS_STEPS] [--tests_dt TESTS_DT]
                               [--tests_entropy_nx TESTS_ENTROPY_NX]
                               [--tests_entropy_xlim TESTS_ENTROPY_XLIM]
                               [--tests_Q_sweep TESTS_Q_SWEEP]
                               [--tests_maxK_sweep TESTS_MAXK_SWEEP]
```

```
[--tests_max_numpy_nmax TESTS_MAX_NUMPY_NMAX]
```

Standalone fast SOE→MPO/TT Landau-Hermite (JAX-first): Fig1 + Fig2 panels.

options:

```
-h, --help           show this help message and exit
--backend {jax,numpy}
--nmax NMAX
--Q Q
--maxK MAXK
--integrator {rk2,ssprk3,rk4}
--linearized {on,off}
                                Include linearized (Maxwellian-background) overlays.
--linearized_method {tangent,matrix}
                                How to compute linearized evolution (default: matrix-
                                free tangent linear).
--progress_chunks PROGRESS_CHUNKS
                                If >0, run JAX time stepping in fixed-size chunks and
                                print progress per chunk.
--quiet
                                Reduce prints.
--use_tt
                                Use optional TT/MPO contraction in NumPy backend
                                (debug/scaling).
--tt_tol TT_TOL
--tt_rmax TT_RMAX
--dt_1sp DT_1SP
--dt_2sp DT_2SP
--tmax_1sp TMAX_1SP
--tmax_2sp TMAX_2SP
--steps_1sp STEPS_1SP
--steps_2sp STEPS_2SP
--fig1_ic {prl_m2,twostream}
                                Fig1 IC family for the 1-species case.
--no_enforce_nonneg_ic
                                Disable IC nonnegativity enforcement on diagnostic
                                grids (not recommended at low nmax).
--amp1 AMP1
                                Fig1 1sp control: for prl_m2 it's the 2nd-moment
                                anisotropy amplitude; for twostream it's the stream
                                separation u in vx/vth.
--dT2 DT2
                                Target temperature excess for the hot species in 2sp
                                ICs (positivity-safe).
--outprefix_fig1 OUTPREFIX FIG1
--outprefix_fig2 OUTPREFIX FIG2
--outprefix OUTPREFIX
                                (compat) Alias for --outprefix_fig1; if
                                --outprefix_fig2 is default, set it to
                                OUTPREFIX+'_Fig2'.
--skip_fig1
--skip_fig2
--fig2_strength FIG2_STRENGTH
                                Scales the positive polynomial distortion used for the
                                Fig2 initial condition (larger => farther from
                                Maxwellian).
--seed SEED
--entropy_nx ENTROPY_NX
```

```

--entropy_xlim ENTROPY_XLIM
--run_tests           Run internal correctness/performance tests and write
                      plots into tests_landau_hermite/.
--tests_outdir TESTS_OUTDIR
--tests_nmax_list TESTS_NMAX_LIST
--tests_reps_rhs TESTS_REPS_RHS
--tests_reps_bench TESTS_REPS_BENCH
--tests_steps TESTS_STEPS
--tests_dt TESTS_DT
--tests_entropy_nx TESTS_ENTROPY_NX
                      Entropy grid resolution used in --run_tests physics
                      checks (smaller is faster).
--tests_entropy_xlim TESTS_ENTROPY_XLIM
                      Entropy grid half-width used in --run_tests physics
                      checks.
--tests_Q_sweep TESTS_Q_SWEEP
                      Comma-separated Q values for convergence sweeps in
                      --run_tests.
--tests_maxK_sweep TESTS_MAXK_SWEEP
                      Comma-separated maxK values for convergence sweeps in
                      --run_tests.
--tests_max_numpy_nmax TESTS_MAX_NUMPY_NMAX
                      In --run_tests, only do NumPy-vs-JAX comparisons up to
                      this nmax (NumPy gets slow for larger nmax).

```

2 Physical model: Landau collision operator

2.1 Continuous operator (context)

For species a colliding with species b , the Landau collision operator can be written (in one common form) as

$$Q_{ab}(f_a, f_b)(\mathbf{v}) = \nabla_{\mathbf{v}} \cdot \int_{\mathbb{R}^3} \mathbf{U}(\mathbf{v} - \mathbf{v}') [f_b(\mathbf{v}') \nabla_{\mathbf{v}} f_a(\mathbf{v}) - f_a(\mathbf{v}) \nabla_{\mathbf{v}'} f_b(\mathbf{v}')] d\mathbf{v}', \quad (1)$$

where \mathbf{U} is the Landau tensor kernel (Coulomb interaction case). This operator is:

- **bilinear** in (f_a, f_b) ,
- **conservative**: it preserves total number, total momentum, and total energy (appropriate combinations across species),
- **entropy producing**: for the *nonlinear* operator, relative entropy to a local Maxwellian is non-increasing in time (H-theorem).

Standard references include Landau (1936), Rosenbluth–MacDonald–Judd (1957), and modern expositions such as Helander & Sigmar (2002) and Villani (2002).

2.2 What the script discretizes

The code discretizes *velocity* in a tensor-product Hermite basis and evolves the corresponding expansion coefficients in time for spatially homogeneous relaxation. All spatial dependence is absent: the ODE is in coefficient space.

3 Normalization, basis, and coefficient tensor representation

3.1 Thermal normalization

Each species has a reference thermal speed v_{th} and mass m . The code uses normalized velocity coordinates

$$x = v_x/v_{\text{th}}, \quad y = v_y/v_{\text{th}}, \quad z = v_z/v_{\text{th}}. \quad (2)$$

The convention used throughout is

$$v_{\text{th}} = \sqrt{\frac{2T_{\text{eq}}}{m}} \iff T_{\text{eq}} = \frac{mv_{\text{th}}^2}{2}, \quad (3)$$

where T_{eq} is the equilibrium temperature corresponding to the chosen normalization.

3.2 Hermite basis used in the code

The script uses a Hermite–Gaussian basis in each dimension:

$$\psi_n(x) = \frac{1}{\sqrt{\pi}} \frac{H_n(x)}{\sqrt{2^n n!}} e^{-x^2}, \quad (4)$$

where H_n is the physicists’ Hermite polynomial. The code implements ψ_n in `psi_1d`. This is a standard choice for Hermite spectral methods in kinetic theory; it provides an efficient ladder structure for derivatives and velocity multiplications.

3.3 Tensor coefficient representation

For a truncation parameter n_{max} , define $p = n_{\text{max}} + 1$. The coefficient tensor is

$$f[\alpha, \beta, \gamma], \quad \alpha, \beta, \gamma \in \{0, \dots, n_{\text{max}}\}, \quad (5)$$

so that

$$f(x, y, z) \approx \sum_{\alpha, \beta, \gamma=0}^{n_{\text{max}}} f[\alpha, \beta, \gamma] \psi_\alpha(x) \psi_\beta(y) \psi_\gamma(z). \quad (6)$$

The code keeps this 3D “cube” representation (shape (p, p, p)) throughout to avoid flattening overhead and to make tensor contractions explicit and JIT-friendly.

3.4 Moment diagnostics and invariants

From the Hermite coefficients, the script extracts:

- density n ,
- momentum \mathbf{P} ,
- total kinetic energy W ,
- temperature $T = \frac{2}{3} \frac{W}{n}$,
- anisotropy measure $A = (T_z - T_x)/T_{\text{avg}}$.

These are computed from low-order coefficients (see `invariants_from_tensor` and `temperature_components_hat`). Conservation of (n, \mathbf{P}, W) (or total invariants for 2 species) is one of the most stringent correctness checks for collision operators.

4 Fast collision evaluation: SOE→MPO/TT

4.1 Why naive evaluation is expensive

In a generic coefficient formulation, the bilinear operator has the form

$$\frac{df_k}{dt} = \sum_{k_a, k_b} C_{k, k_a, k_b}^{ab} f_{a, k_a} f_{b, k_b}, \quad (7)$$

which is $O(N^3)$ if the dense tensor C^{ab} is formed and contracted directly (with $N = p^3$). Even storing C^{ab} is prohibitive.

The script avoids forming C^{ab} and instead uses a structured evaluation that exploits separability of the underlying integrals and the tensor-product basis.

4.2 SOE / separable quadrature idea

The coefficient formulas in this method contain factors of the form $\frac{1}{2A+1}$ (with A an integer index arising from the Coulomb kernel moments). The key separability identity is:

$$\frac{1}{2A+1} = \int_0^1 s^{2A} ds. \quad (8)$$

Approximating the integral using Gauss–Legendre quadrature on $[0, 1]$ gives:

$$\frac{1}{2A+1} \approx \sum_{q=1}^Q w_q s_q^{2A}, \quad (9)$$

which is a *sum of separable terms* across dimensions because A itself decomposes into a sum of 1D indices. In the code, `leggauss_01_np` constructs (s_q, w_q) , and the resulting q -terms are handled via a small loop over q (either explicit loop in NumPy or vectorized/JAX loop).

This “SOE” terminology is used in the code base to mean “sum of separable terms” (often also called separated representations, sum-of-products, etc.). Related ideas are common in fast algorithms for integral operators (e.g. Beylkin & Monzón; Hackbusch & Khoromskij).

4.3 Separable Hankel structure

The algorithm also uses a Hankel-like dependence on combined indices $K = k'_p + k_b + \text{shift}$. In the code, this enters through the table

$$F_q[n] = \begin{cases} 0, & n \text{ odd}, \\ (-1)^{n/2} t_{n/2} s_q^n, & n \text{ even}, \end{cases} \quad (10)$$

where

$$t_n = \frac{\sqrt{(2n)!}}{2^n n!}. \quad (11)$$

See `t_factor_np` and `Fq_table_np`. The truncation parameter `maxK` limits the maximum Hankel index needed; the docstring in `build_model_tables_np` explains the rule of thumb (`maxK` comfortably above $\sim 3n_{\max}$).

4.4 MPO viewpoint: tensor-product operator contraction

After SOE separation, the operator can be evaluated through a sequence of *1D* mode products along the x , y , and z axes. In quantum many-body language this is a matrix-product operator (MPO); in numerical linear algebra it is a Kronecker-structured operator.

Concretely, the code precomputes (for each quadrature node q , tensor component indices i, j , and term type) three 1D matrices that act along each dimension. Applying the operator to a 3D tensor then reduces to three successive contractions (a Kronecker product action), implemented via `einsum/tensordot` and fused by XLA in the JAX path.

4.5 TT (tensor train) compression

For larger n_{\max} , some intermediate tensors can become costly. A tensor-train (TT) factorization approximates a 3D tensor $X \in \mathbb{R}^{p \times p \times p}$ by low-rank factors (TT-SVD). The script includes an optional TT rounding path (enabled via `--use_tt` in the NumPy backend) to reduce intermediate ranks and memory traffic when exploring higher n_{\max} .

TT/MPS references: Oseledets (2011), Hackbusch (2012), and related MPO/MPS reviews (Schollwöck 2011; Orús 2014).

4.6 Algorithmic walkthrough (as implemented)

This subsection summarizes the concrete data flow in `landau_hermite_jax.py`.

Inputs and shapes. Fix a truncation n_{\max} and set $p = n_{\max} + 1$. Each distribution is a tensor $f \in \mathbb{R}^{p \times p \times p}$ in the Hermite basis.

Step 0: precompute tables (once per $(n_{\max}, Q, \text{maxK})$ and species pair). The function `build_model_tables_np` builds a container of dense arrays:

- Gauss-Legendre nodes/weights (s_q, w_q) on $[0, 1]$ for the separable quadrature.
- 1D mixing coefficients `a1d` encoding relative/COM transforms with the species-dependent angle ($\cos \theta, \sin \theta$).
- A dense 1D coefficient table `P1D` that maps intermediate indices back to Hermite indices (a 1D “dual \leftrightarrow primal” transform used repeatedly in each dimension).
- The separable Hankel tables `Fq[q, K]` which implement the even- K Coulomb moment factors multiplied by s_q^K (see `Fq_table_np`).
- Preassembled 1D matrices `M_buildS[q, i, j, term, dim]` used to build auxiliary tensors S_1^{ij} and S_2^{ij} from the background distribution (next paragraph).

In the JAX backend these arrays are transferred to device once and treated as constants by JIT.

Step 1: build auxiliary tensors from f_b . For a cross-collision $Q_{ab}(f_a, f_b)$, the algorithm first constructs a small set of intermediate tensors (denoted S_1^{ij} and S_2^{ij} in the code) that play the role of Rosenbluth-potential moments in this Hermite formulation. Each S_ℓ^{ij} has shape (p', p', p') where $p' = 2n_{\max} + 2$ (the intermediate index range required by the Hankel structure). The construction is a sum over quadrature nodes q of Kronecker-structured applications:

$$S_\ell^{ij} \approx \sum_{q=1}^Q w_q \left(M_{q,i,j,\ell}^{(x)} \otimes M_{q,i,j,\ell}^{(y)} \otimes M_{q,i,j,\ell}^{(z)} \right) f_b, \quad (12)$$

implemented via three successive 1D contractions. This avoids forming any dense 6D objects.

Step 2: apply S to f_a to produce the RHS. Given S_1, S_2 (from f_b), the RHS for f_a is assembled by applying another Kronecker-structured mapping and then projecting back to the Hermite coefficient cube using the 1D P1D table along each axis. This is the “MPO” viewpoint: the full operator factorizes over dimensions, so applying it is a sequence of mode products rather than a giant tensor contraction.

Optional: TT compression (NumPy path). When `--use_tt` is enabled in the NumPy backend, selected intermediate tensors may be TT-rounded to reduce ranks. The JAX backend currently relies primarily on XLA fusion rather than TT-SVD, because SVD-based compression is relatively expensive on CPU for the small-to-moderate p used in the main figures.

Complexity. The dominant cost scales roughly like $O(Q p^4)$ for the structured contractions (with a small constant), rather than $O(p^9)$ for a dense bilinear tensor.

5 Nonlinear vs linearized evolution

5.1 Nonlinear operator

The main ODE uses the full bilinear operator:

$$1\text{sp: } \frac{df}{dt} = Q_{11}(f, f), \quad 2\text{sp: } \frac{df_a}{dt} = Q_{ab}(f_a, f_b), \quad \frac{df_b}{dt} = Q_{ba}(f_b, f_a). \quad (13)$$

5.2 Linearization about a Maxwellian

Let M denote a Maxwellian equilibrium (in the *fixed* normalized coordinates for each species). For a perturbation h with $f = M + h$ and $\|h\| \ll 1$,

$$Q(M + h, M + h) = Q(M, M) + Q(h, M) + Q(M, h) + O(h^2). \quad (14)$$

Since $Q(M, M) = 0$, the linearized operator is

$$L(h) = Q(h, M) + Q(M, h). \quad (15)$$

For two species, the linearized Jacobian couples δf_a and δf_b through the mixed terms.

5.3 Entropy vs free energy for linearized dynamics

The nonlinear Landau operator satisfies an H-theorem: relative entropy to the *instantaneous* local Maxwellian decreases. The script measures this via a grid-based KL divergence:

$$\mathcal{D}(t) = \int f \log\left(\frac{f}{M_{\text{local}}}\right) dv. \quad (16)$$

However, this KL functional is *not* expected to be monotone for tangent-linear evolution (and can become ill-defined if the linearized solution produces slight negative values due to truncation).

For linearized evolution, a standard Lyapunov functional is the quadratic “free energy” (second variation of entropy at a Maxwellian):

$$\mathcal{F}(t; M) = \int \frac{(f - M)^2}{M} dv. \quad (17)$$

The script therefore plots $\mathcal{D}/\mathcal{D}_0$ for nonlinear curves and $\mathcal{F}/\mathcal{F}_0$ for linearized curves in the main panels, and it tests monotonicity for each in `--run_tests`.

6 Time stepping and JAX implementation

6.1 Integrator

The script uses explicit SSPRK3 (Shu–Osher) by default, with options for RK2 and RK4. SSPRK3 offers a good stability/accuracy compromise for dissipative problems while using only 3 RHS evaluations per step. See Shu & Osher (1988) and Gottlieb–Shu–Tadmor (2001).

6.2 JAX performance strategy

Key performance choices:

- Enable float64 (`jax_enable_x64`) for accurate invariants and Maxwellian fixed-point checks.
- Treat all precomputed tables as constants to JIT (single device transfer).
- Use `lax.scan` for the time loop so the full stepper is compiled once.
- Keep shapes static (fixed $p = n_{\max} + 1$ and fixed Q per run).
- Avoid Python loops over tensor indices; rely on vectorized contractions.

These are typical “best practices” for XLA-based array programming systems.

7 Initial conditions used in Fig. 1 and Fig. 2

7.1 Fig. 1: one-species IC (two-stream-like, positivity-safe)

The default Fig. 1 one-species IC is designed to be qualitatively “two-stream-like” without any negative values. It is constructed as an even mixture of two shifted Maxwellians in v_x :

$$f(x, y, z) \propto \frac{1}{2} \left(M(x - u) + M(x + u) \right) M(y) M(z), \quad (18)$$

with u set by `--amp1` (interpreted as stream separation for `--fig1_ic_twostream`). The code projects the 1D x dependence onto the truncated $\{\psi_n\}$ basis and optionally scales the non-Maxwellian component to enforce nonnegativity on diagnostic grids.

7.2 Fig. 1: two-species IC (positivity-safe hot/equilibrium)

Representing an extremely cold Maxwellian at low n_{\max} in a fixed normalized coordinate can lead to Gibbs-like oscillations and negative reconstructed f . To avoid this, the two-species IC used in this standalone script heats one species using a positive isotropic polynomial distortion of the equilibrium Maxwellian (exactly representable at low order), while keeping the other species near equilibrium. The parameter `--dT2` sets the target temperature excess for the hot species.

7.3 Fig. 2: far-from-Maxwellian ICs

Fig. 2 uses a positive polynomial distortion (including quartic terms when $n_{\max} \geq 4$) to create strongly non-Maxwellian shapes while maintaining nonnegativity on diagnostic grids. The strength is controlled by `--fig2_strength`. The two-species Fig. 2 IC also tunes the temperature separation in a positivity-safe way.

Table 1: Test sweep summary (JAX backend). Reported times are steady-state. Conservation rates are for cross-collision totals $d/dt(n, P, W)$ evaluated on a random near-Maxwellian state.

n_{\max}	$N = (n_{\max}+1)^3$	t_{self} [ms]	t_{cross} [ms]	t_{int} [s]	$\ RHS(M)\ /\ M\ $ (cross)	$ dW_{\text{tot}}/dt $
2	27	0.093	0.177	0.014	4.23e-16	1.63e-15
3	64	0.280	0.556	0.044	4.35e-16	8.50e-12
4	125	0.584	1.141	0.092	5.98e-16	4.78e-11
5	216	0.854	1.711	0.133	6.00e-16	3.26e-10

8 Figures produced by the script

8.1 Main panels

8.2 Test suite summary and plots

The `--run_tests` mode runs a sweep over n_{\max} , checks Maxwellian fixed points, cross-invariant rates, finite-difference consistency of the linearization (when feasible), and short-run performance. It also performs physics-style checks of monotonicity (nonlinear KL entropy and linear quadratic free energy) and generates convergence sweeps in (Q, maxK) .

Interpreting the test outputs. The tests are designed around standard invariants and stability properties of the Landau operator:

- **Maxwellian fixed point:** $RHS(M) \approx 0$ is a strict algebraic property of the collision operator. In floating point, the residual should be near roundoff if the discretization/tables are correct.
- **Conservation:** for cross-collisions, Q_{ab} and Q_{ba} exchange momentum/energy but must conserve totals. The test plots $|d/dt(n_{\text{tot}}, P_{\text{tot}}, W_{\text{tot}})|$ computed directly from the RHS.
- **H-theorem diagnostics:** nonlinear KL entropy to the instantaneous local Maxwellian should be non-increasing (up to truncation/time-discretization). Linearized dynamics should instead be monitored with quadratic free energy about the fixed Maxwellian background.
- **Linearization correctness:** two complementary checks are used: (i) finite-difference self-consistency (NumPy, small n_{\max}) and (ii) small-perturbation agreement between nonlinear and tangent-linear time traces (JAX).
- **Performance:** steady-state RHS timing (after compilation) and a short integration benchmark.




Figure 1: Fig. 1 panel produced by `landau_hermite_jax.py`.




Figure 2: Fig. 2 panel produced by `landau_hermite_jax.py`.




Figure 3: Per-RHS-call runtime scaling across n_{\max} (from the latest test run).




Figure 4: Maxwellian fixed-point residual $\|RHS(M)\|/\|M\|$ across n_{\max} .




Figure 5: Cross-collision conservation-rate check: total invariant rates computed from the RHS on a random near-Maxwellian state.




Figure 6: Physics-style monotonicity check: nonlinear KL entropy to instantaneous local Maxwellians (solid) and linearized quadratic free energy about fixed Maxwellians (dashed).




Figure 7: Backend consistency between NumPy and JAX for small n_{\max} where NumPy comparisons are enabled.




Figure 8: Finite-difference self-consistency check of the linearization (NumPy, where enabled).

1sp: nonlinear vs tangent-linear (small perturbation)




Figure 9: Small-perturbation check: nonlinear vs tangent-linear for the 1-species problem.

2sp: dT nonlinear vs tangent-linear (small perturbation)




Figure 10: Small-perturbation check: nonlinear vs tangent-linear for the 2-species temperature-exchange metric.




Figure 11: JAX compile-time scaling across n_{\max} (first call).




Figure 12: Short integration benchmark (NumPy vs JAX where enabled).

9 Interpreting results and limitations

9.1 Positivity

Hermite spectral truncations do *not* guarantee pointwise positivity of the reconstructed distribution function, even if the underlying physical solution is nonnegative. This is a well-known spectral/Gibbs phenomenon. To keep the plotted slices physically meaningful, the default Fig. 1 one-species IC is built from a manifestly nonnegative two-stream mixture and is optionally “damped” in high modes to remain nonnegative on diagnostic grids.

9.2 Monotonicity

For the nonlinear operator, $\mathcal{D}(t)$ (KL divergence to the instantaneous local Maxwellian) is the referee-proof diagnostic. For the linearized evolution, the correct monotone quantity is the quadratic free energy $\mathcal{F}(t; M)$ about the fixed Maxwellian background, not KL to a time-dependent Maxwellian.

9.3 Truncation and parameter choices

The separation accuracy is controlled by:

- n_{\max} : Hermite resolution (spectral truncation),
- Q : quadrature nodes in the SOE separation,
- \maxK : maximum Hankel index in the Coulomb moment tables.

The test suite includes Q/\maxK sweeps and Maxwellian fixed-point residuals to help select safe defaults.

10 References (selected)

References

- [1] L. D. Landau, “Kinetic equation for the Coulomb effect,” *Physikalische Zeitschrift der Sowjetunion* (1936).
- [2] M. N. Rosenbluth, W. M. MacDonald, and D. L. Judd, “Fokker–Planck equation for an inverse-square force,” *Physical Review* **107** (1957).
- [3] P. Helander and D. J. Sigmar, *Collisional Transport in Magnetized Plasmas*, Cambridge Univ. Press (2002).
- [4] C. Villani, “A review of mathematical topics in collisional kinetic theory,” in *Handbook of Mathematical Fluid Dynamics*, Vol. I (2002).
- [5] C.-W. Shu and S. Osher, “Efficient implementation of essentially non-oscillatory shock-capturing schemes,” *Journal of Computational Physics* **77** (1988).
- [6] S. Gottlieb, C.-W. Shu, and E. Tadmor, “Strong stability-preserving high-order time discretization methods,” *SIAM Review* **43** (2001).
- [7] I. V. Oseledets, “Tensor-train decomposition,” *SIAM Journal on Scientific Computing* **33** (2011).
- [8] W. Hackbusch, *Tensor Spaces and Numerical Tensor Calculus*, Springer (2012).

- [9] U. Schollwöck, “The density-matrix renormalization group in the age of matrix product states,” *Annals of Physics* **326** (2011).
- [10] R. Orús, “A practical introduction to tensor networks: Matrix product states and projected entangled pair states,” *Annals of Physics* **349** (2014).
- [11] G. H. Golub and J. H. Welsch, “Calculation of Gauss quadrature rules,” *Mathematics of Computation* **23** (1969).
- [12] J. P. Boyd, *Chebyshev and Fourier Spectral Methods*, 2nd ed., Dover (2001). (Background on spectral truncation/oscillations.)
- [13] H. Grad, “On the kinetic theory of rarefied gases,” *Communications on Pure and Applied Mathematics* **2** (1949).
- [14] G. Beylkin and L. Monzón, “Approximation by exponential sums revisited,” *Applied and Computational Harmonic Analysis* (2005).
- [15] J. Bradbury et al., “JAX: composable transformations of Python+NumPy programs,” (2018).
<https://github.com/jax-ml/jax>